

Assignment 3 OS

Ansh Prakash
2016CS10367

Himanshu
2016CS10348

IMPLEMENTATION

Each process is have a cid in its struct proc, which tells to which container it belongs.

Our Implementation is most in user space with some tweaks in the kernel space.

Create_container:

This initializes the data structure required for the implementation

Leave_container:

This simply free the process from the container by assigning it the cid = -1

Join_container:

This change the cwd of the process to container dir and provide it a slot in container process table

Destroy_Container:

Kills all the process running in it. Empty the container folder(i.e delete all the files for the container)

code in kernel space

```
int
create_container(int cid){
    struct container container;
    if (no_of_containers == 0){
        // Initialising the Data Structure for the first time(assume we never delete a container(which is
        not true))
        for (int i = 0; i < MAX_CONTAINER; ++i){
            struct container temp;
            temp.cid = -1;
            containers[i] = temp;
        }
    }
    if (no_of_containers+1 == MAX_CONTAINER){
        return(-1);
    }
    no_of_containers += 1;
    containers[cid] = container;
    container.cid = cid;
    container.current_running_process = -1;
    container.next_pid = 0;
    for(int i=0; i<MAX_NPROC; i++){
        struct proc p;
        p.pid = -1;
```

```

    container.ptable.proc[i] = p;
}
cprintf("Created Container %d \n",cid);
return((uint)cid);
}

```

DESTROY_CONTAINER

```

uint
destroy_container ( uint container_id )
{
    containers[container_id].cid = -1;
    acquire(&helperlock.lock);
    struct proc *p;
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->cid == container_id){
            kill(p->pid);
        }
    }
    release(&helperlock.lock);
    cprintf("Destroy Container %d \n",container_id);
    return(container_id);
}

```

JOIN_CONTAINER

```

uint
join_container(uint container_id)
{
    // struct ctable p = containers[container_id].ptable;
    struct proc *p = myproc();
    if(p->cid != -1 || containers[container_id].cid == -1){
        return(-1);
    }
    p->cid = container_id;
    // p->state = WAITING;
    if (containers[container_id].next_pid == MAX_NPROC-1){
        return(-1);
    }
    containers[container_id].ptable.proc[containers[container_id].next_pid] = *p;
    containers[container_id].next_pid +=1;
    cprintf("Joined Container %d ",container_id);
    cprintf("my pid is  %d \n",p->pid);
    return(0);
}

```

Leave Container

```

uint
leave_container(void)
{
    struct proc *p = myproc();
    int cid = p->cid;
}

```

```

p->cid = -1;
cprintf("Leave Container %d \n",cid);
return(0);
}

```

PS:

This is Implemented in Kernel space

```

void
ps()
{
    struct proc *mp = myproc();
    int mycid = mp->cid;
    struct proc *p;
    acquire(&ptable.lock);
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->state != UNUSED && p->cid == mycid){
            cprintf("pid:%d name:%s\n",p->pid,p->name);
        }
    }
    release(&ptable.lock);
}

```

Helper sys call to get cid in user space

```

int
getcid(){
    struct proc *mp = myproc();
    int mycid = mp->cid;
    return(mycid);
}

```

PROC STRUCT:

```

// Per-process state
struct proc {
    uint sz;                // Size of process memory (bytes)
    pde_t* pgdir;           // Page table
    char *kstack;           // Bottom of kernel stack for this process
    enum procstate state;    // Process state
    int pid;                 // Process ID
    struct proc *parent;     // Parent process
    struct trapframe *tf;    // Trap frame for current syscall
    struct context *context; // swtch() here to run process

```

```

void *chan;           // If non-zero, sleeping on chan
int killed;           // If non-zero, have been killed
struct file *ofile[NOFILE]; // Open files
struct inode *cwd;     // Current directory
char name[16];         // Process name (debugging)
int cid;
};

```

CONTAINER STRUCT

```

struct container {
    int cid;
    int next_pid;
    int current_running_process;
    struct cptable ptable;
};

```

USER SPACE

CODE

ls and COW-mechanism have been implemented in USER space.

```

#include "types.h"
#include "stat.h"
#include "user.h"
#include "fs.h"
#include "fcntl.h"

```

```

//convert int to string
void
itoa(int num , char* ret)
{
    int rem, len = 0;
    int n = num;
    while(n!=0){
        len++;
        n = n/10;
    }
    for(int i=0;i<len;i++){
        rem = num%10;
        num = num/10;
        ret[len - (i+1)] = rem + '0';
    }
    ret[len] = '\0';
}

```

```

char*
my_fmtname(char *path)
{
    static char buf[DIRSIZ+1];
    char *p;

    // Find first character after last slash.
    for(p=path+strlen(path); p >= path && *p != '/'; p--)
        ;
    p++;

```

// Return blank-padded name.

```

if(strlen(p) >= DIRSIZ)
    return p;
memmove(buf, p, strlen(p));
memset(buf+strlen(p), ' ', DIRSIZ-strlen(p));
return buf;
}

void
my_ls(char *path)
{
    char buf[512], *p;
    int fd;
    struct dirent de;
    struct stat st;

    if((fd = open(path, 0)) < 0){
        printf(2, "ls: cannot open %s\n", path);
        return;
    }

    if(fstat(fd, &st) < 0){
        printf(2, "ls: cannot stat %s\n", path);
        close(fd);
        return;
    }

    switch(st.type){
    case T_FILE:
        printf(1, "%s\n", my_fmtname(path));
        break;

    case T_DIR:
        if(strlen(path) + 1 + DIRSIZ + 1 > sizeof buf){
            printf(1, "ls: path too long\n");
            break;
        }
        strcpy(buf, path);
        p = buf+strlen(buf);
        *p++ = '/';
        while(read(fd, &de, sizeof(de)) == sizeof(de)){
            if(de.inum == 0)
                continue;
            memmove(p, de.name, DIRSIZ);
            p[DIRSIZ] = 0;
            if(stat(buf, &st) < 0){
                printf(1, "ls: cannot stat %s\n", buf);
                continue;
            }
            printf(1, "%s\n", my_fmtname(buf));
        }
        break;
    }
    close(fd);
}

void
destroy_container_wrapper(int cid)
{
    char path[4];
    path[0] = 'c';
    char conv[3];
    itoa(cid, conv);
    path[1] = conv[0];

```

```
path[2] = conv[1];
path[3] = '\0';
```

```
char deldir[7];
deldir[0] = '.';
deldir[1] = '.';
deldir[2] = '/';
for(int i=0 ;i<4;i++){
    deldir[3+i] = path[i];
}
```

```
unlink(deldir);
mkdir(path);
destroy_container(cid);
}
```

```
void
cont_ls()
{
    my_ls(".");
}
```

```
void
init_containers_resources(){
    //setting up the file system
    mkdir("c1");
    mkdir("c2");
    mkdir("c3");
    mkdir("c4");
    mkdir("c5");
    mkdir("c6");
    mkdir("c7");
    mkdir("c8");
    mkdir("c9");
    mkdir("c10");
    mkdir("c11");
    mkdir("c12");
    mkdir("c13");
    mkdir("c14");
    mkdir("c15");
    mkdir("c16");
    mkdir("c17");
    mkdir("c18");
    mkdir("c19");
    mkdir("c20");
}
```

```
void
join_container_wrapper(int i){
    char cid_s[4];
    cid_s[0] = 'c';
    char conv[3];
    itoa(i,conv);
    cid_s[1] = conv[0];
    cid_s[2] = conv[1];
    cid_s[3] = '\0';
    chdir(cid_s);
    join_container(i);
}
```

```

int
leave_container_wrapper()
{
    chdir("/");
    return(leave_container());
}

void
create_file(char* filename)
{
    int fd;
    if((fd = open(filename,O_CREATE))<0){
        printf(1," not able to create the file\n");
    }
    close(fd);
}

void
copyfile(char* file1,char* file2)
{
    int fd0,fd1,n;
    char buf[512];
    if((fd0 = open(file1,O_RDONLY))<0){
        printf(1," error while copying \n");
    }
    create_file(file2);
    if((fd1 = open(file1,O_RDWR))<0){
        printf(1," error while copying \n");
    }
    while((n = read(fd0,buf,sizeof(buf))) > 0 ){
        write(fd1,buf,n);
    }
    close(fd0);
    close(fd1);
}

/*
 * mode : ronly 0
 *       : wronly 1
 *       : rdwr 2
 */
int
open_file(char* filename,int mode)
{
    int fd = -1;
    if(mode == 0 ){
        fd = open(filename,O_RDONLY);
    }
    else if(mode == 1){
        fd = open(filename,O_WRONLY);
    }
    else if(mode == 2){
        fd = open(filename,O_RDWR);
    }
}

if(fd<0){
    char path[30];
    path[0] = '/';
    int i = 0;
    while(filename[i]!='\0'){
        path[i+1] = filename[i];
        i++;
    }
}

```

```

}
if((fd = open(path,O_RDWR))<0){
    printf(1," no such file exist\n");
    return(-1);
}
int cid = getcid();
char path_t[40];
path_t[0] = '/';
path_t[1] = 'c';
char conv[3];
itoa(cid,conv);
int j = 2;
while(conv[j-2]!='\0'){
    path_t[j] = conv[j-2];
    printf(0,"%s\n", conv);
    j++;
}
path_t[j] = '/';
j++;
i = 0;
while(filename[i]!='\0'){
    path_t[i+j] = filename[i];
    i++;
}
path_t[i+j] = '\0';
if(mode == 0){
    fd = open(path,O_RDONLY);
}
else if(mode == 1){
    if((fd = open(path,O_WRONLY))<0){
        printf(1," no such file exist\n");
    }
    close(fd);
    copyfile(path,path_t);
    fd = open(filename,O_WRONLY);
    return(fd);
}
else if(mode == 2){
    if((fd = open(path,O_RDWR))<0){
        printf(1," no such file exist\n");
    }
    close(fd);
    printf(0,"file apth_t %s\n",path_t);
    copyfile(path,path_t);

    fd = open(filename,O_RDWR);
    return(fd);
}
if(fd<0){
    printf(1," no such file exist\n");
}
}
return(fd);
}

```

```

void
cont_cat(char* file)
{
    int fd,n;
    char buf[512];
    if((fd = open_file(file,0))<0){
        printf(1," error while copying \n");
    }
}

```



```

}

while((n = read(fd,buf,sizeof(buf))) > 0 ){
    printf(0,"%s",buf);
}
close(fd);
}

int main(void)
{

    init_containers_resources();
        create_container(1);
    create_container(2);
    create_container(3);
    //////////////////////////////////////
    int pid = fork();
    if(pid == 0){
        join_container_wrapper(1); // called only by child created by preceeding fork call .
        create_file("hello");
        int fd;
        fd =open_file("hello",2);
        char raw[5] = "ansh\n";
        int n =5;
        write(fd,raw,n);
        close(fd);
        printf(0,"Testing ls \n\n");
        cont_ls();
        // printf(0,"Testing ps \n\n\n");
        // ps();
        exit();
    }
    // destroy_container_wrapper(1);

    pid = fork();
    if(pid == 0){
        join_container_wrapper(2); // called only by child created by preceeding fork call .
        create_file("YOYO");
        int fd;
        fd =open_file("YOYO",2);
        char raw[5] = "five\n";
        int n =5;
        write(fd,raw,n);
        close(fd);
        printf(0,"Testing ls \n\n\n");
        cont_ls();
        printf(0,"Testing ps \n\n\n");
        ps();
        exit();
    }
    //////////////////////////////////////

    //////////////////////////////////////COPY ON WRITE TESTING //////////////////////////////////////
    create_file("ROOT");
    int fdt;
    fdt =open_file("ROOT",2);
    char raw[5] = "root\n";
    int n =5;
    write(fdt,raw,n);
    close(fdt);

    int pid = fork();
    if(pid == 0){

```

```

join_container_wrapper(2); // called only by child created by preceeding fork call .
int fd;
cont_cat("ROOT");//This will read from the root file
fd =open_file("ROOT",2);
char raw[9] = "anshroot\n";
int n =9;
write(fd,raw,n);
close(fd);
printf(0,"Testing ls \n\n");
cont_ls();
cont_cat("ROOT");
exit();
}
//////////COPY ON WRITE TESTING //////////
wait();
exit();
}

```

FILE ISOLATION STRATEDY:

For resource isolation the user spcae(create_container.c) code first create folders for each of the possible container and then any file created in that kernel will be stored in that folder.

This implementation keep in mind that the container process can see the root dir.

Special functionality

```
void copyfile(char* file1,char* file2);
```

This function copy the content of file1 to file 2.
